Project RIFT: From Reflexes to Geometric Anticipation

A Physical Approach to Intelligence under Irreversibility

Akshay John
Independent Researcher

January 2026

Abstract

Modern artificial intelligence systems rely heavily on prediction, probabilistic inference, and reward optimization. These approaches implicitly assume reversibility, recoverability, and tolerance for delayed action. In contrast, many real-world survival scenarios impose irreversible consequences, metabolic costs, and sparse delayed outcomes.

This paper presents Project RIFT, a systematic experimental investigation into intelligence under irreversible constraints. Across more than fifty controlled simulations, we study agents operating in environments where survival requires early, irreversible commitment based on faint causal cues. Actions incur continuous metabolic cost, and success or failure is resolved only after a delayed hazard.

We demonstrate a consistent failure mode across predictive agents—including Transformers, recurrent neural networks, and reinforcement learners. These agents hedge under uncertainty, delaying commitment until it is no longer survivable. In contrast, simple geometric and morphological agents succeed by collapsing uncertainty through irreversible physical commitment rather than internal prediction.

A final controlled experiment (RIFT-A) isolates this dissociation: a non-learning geometric threshold agent achieves nontrivial survival, while all predictive agents fail despite access to identical observations. These results suggest that anticipation can emerge as a property of physical and geometric constraints under irreversibility, and that prediction-based intelligence is structurally misaligned with certain classes of survival problems.

Contents

1 Introduction **4**

2 Core Hypothesis and Experimental Scope **4**
2.1 Experimental Scope 4
2.2 Central Hypothesis 4

3 Experimental Framework **5**

4 Phase I: Reflex Without Time (Ballistic Experiments) **5**
4.1 $\texttt{rift}_b allistic.py - - - BaseReflex$ 5
4.2 $\texttt{rift}_b allistic_1.py - - - LethalVariant$ 6

5 Phase II: Mortality and Fear **6**
5.1 Objective 6
5.2 `rift_scare.py` — The Mortal Body 7

5.2.1 Concept . . . 7
5.2.2 Environment Setup . . . 7
5.2.3 Key Dynamics . . . 7
5.2.4 Observed Behavior . . . 7
5.2.5 Result: Paranoia . . . 8
5.2.6 Key Insight . . . 8
5.3 Phase II Conclusion . . . 8
5.4 Transition to Phase III . . . 8

6 Phase III: Energy, Cost, and Constraint 8
6.1 Objective . . . 8
6.2 `rift_ischemia.py` — Fatigue of Readiness . . . 9
6.3 `rift_adrenaline.py` — Finite Stamina . . . 9
6.4 `rift_burden.py` — Evolutionary Weight . . . 10
6.5 Phase III Summary . . . 10

7 Phase IV: Sensory Degradation, Ambiguity, and Cognitive Conflict 10
7.1 Objective . . . 10
7.2 `rift_world.py` — Unified Physiological World . . . 11
7.3 `rift_ambiguity.py` — Probabilistic Futures . . . 11
7.4 `rift_dissonance.py` — Cognitive Dissonance . . . 11
7.5 Phase IV Summary . . . 12

8 Phase V: Motor Control, Discontinuity, and the First Breakthrough 12
8.1 Objective . . . 12
8.2 `rift_kinesis.py` — Motor Babbling . . . 13
8.3 `rift_kinesis2.py` — Impedance Control . . . 13
8.4 `rift_contact.py` — Gating . . . 13
8.5 `rift_impulse.py` — The Jump (BREAKTHROUGH) . . . 13
8.6 `rift_switch.py` — Plasticity and Melting . . . 14
8.7 Phase V Summary . . . 14

9 Phase VI: Memory Without Storage — Traces, Engrams, and Physical Recall 14
9.1 Objective . . . 14
9.2 `rift_mnemos.py` — Memory Capacity Test . . . 15
9.3 Iterative Memory Experiments (`rift_mnemos` Series) . . . 15
9.4 The Prediction Trap (`rift_mnemos_p` Series) . . . 15
9.5 `rift_engram.py` — Neural Memory Injection . . . 15
9.6 `rift_phase_transition.py` — Thermodynamic Lock (BREAKTHROUGH) . . . 16
9.7 `rift_homeostasis.py` — Regulation . . . 16
9.8 Phase VI Summary . . . 16

10 Phase VII: Morphology and Geometric Commitment (THE BREAKTHROUGH) 17
10.1 Objective . . . 17
10.2 `rift_morphology.py` — Center-of-Mass Shift . . . 17
10.3 `rift_morphology_v2.py` — Tuned Morphology . . . 17
10.4 Formal Insight: The RIFT Pipeline . . . 18
10.5 Constraint Against Hedging . . . 18

10.6 Comparison With Neural Systems . . . 18
10.7 Conceptual Breakthrough . . . 18
10.8 Phase VII Summary . . . 19

11 Phase VIII: Ecology, Speciation, and Selection Pressure **19**
11.1 Objective . . . 19
11.2 The Failure of Soft Constraints (`rift_ecology` Series) . . . 19
11.3 `rift_min_speciation.py` — Hard Logical Exclusion . . . 20
11.4 `rift_necessity.py` — Contextual Pressure . . . 20
11.5 Pathologies of Adaptation . . . 20
11.6 Phase VIII Summary . . . 20

12 Phase IX: Validation, Falsification, and Collapse Tests **21**
12.1 Objective . . . 21
12.2 `rift_validation.py` — The Inquisitor . . . 21
12.2.1 Test 1: Control (Baseline) . . . 21
12.2.2 Test 2: Stretch (Temporal Shift) . . . 21
12.2.3 Test 3: Noise Injection . . . 21
12.2.4 Test 4: Lobotomy (Ablation) . . . 22
12.3 Statistical Interpretation . . . 22
12.4 Phase IX Summary . . . 22

13 Final Conclusion (PASS 1) **22**
13.1 Environment Design: RIFT-A Commitment World . . . 23
13.2 Agent Architectures Evaluated (Jules Runs) . . . 23
13.3 Quantitative Results and Failure Modes . . . 24
13.4 Failure Analysis: Hedging vs Commitment . . . 25

14 Discussion **25**
14.1 Why Predictive Intelligence Fails . . . 25
14.2 Why Geometric Systems Succeed . . . 25
14.3 Relation to Existing AI Paradigms . . . 26
14.4 Relation to Active Inference and Control-Theoretic Models . . . 26

15 Limitations **26**

16 Conclusion and Future Work **27**
16.1 What PASS 1 Establishes . . . 27
16.2 PASS 2 Directions . . . 27

1 Introduction

Contemporary artificial intelligence systems overwhelmingly rely on predictive modeling, reward optimization, and large-scale parameterized function approximators. While these approaches have yielded impressive results in language modeling and pattern recognition, their applicability to embodied intelligence, survival, and anticipation under irreversible physical constraints remains poorly understood.

Project RIFT was initiated to answer a deceptively simple question:

> *Can intelligence emerge without prediction, rewards, or symbolic memory, purely from physical necessity?*

Rather than designing agents to maximize reward, we construct environments where survival itself imposes irreversible costs, delayed consequences, and metabolic trade-offs. Agents are not trained to be intelligent; instead, intelligence is made unavoidable.

This paper documents the complete experimental journey of Project RIFT, from early reflex-based failures to the discovery of geometric anticipation.

2 Core Hypothesis and Experimental Scope

Project RIFT is guided by an experimentally grounded hypothesis rather than an architectural prescription. The project does not assume that prediction, learning, or neural representations are inherently flawed. Instead, it asks whether these mechanisms remain effective under irreversible physical constraints.

2.1 Experimental Scope

RIFT focuses on environments characterized by:

- Irreversible actions with no opportunity for rollback
- Continuous metabolic cost for maintaining readiness
- Sparse and delayed resolution of success or failure
- Catastrophic consequences for delayed commitment

Within this scope, agents were evaluated using a range of architectures, including reflexive controllers, neural predictors, reinforcement learners, and geometric mechanisms.

2.2 Central Hypothesis

> *In environments with irreversible cost and delayed consequence, intelligence emerges through early physical commitment rather than predictive optimization.*

This hypothesis yields several testable implications:

- Actions must sometimes be irreversible to enable anticipation
- Maintaining readiness must incur real metabolic cost
- Hesitation and probabilistic hedging should be penalized by the environment
- Memory must be physically grounded to resist drift and hallucination

All subsequent phases of Project RIFT are designed to test these implications empirically.

3 Experimental Framework

All RIFT experiments share a common physical substrate:

- A 2D spatial grid (typically 32x32 or 48x48)
- Multiple continuous fields (stress, velocity, trace, health)
- Local interactions governed by diffusion, damping, and external forces

Agents interact with the world through minimal scalar or geometric actions (e.g., brace, impulse, posture shift). Neural components, when present, are deliberately constrained.

Terminal outputs, diagnostics, and failures were logged and preserved (see `Terminal Output.md`).

4 Phase I: Reflex Without Time (Ballistic Experiments)

Objective. The goal of Phase I was to establish the simplest possible baseline: can a purely reactive system survive a delayed lethal event using only a scalar reflex?

This phase deliberately excluded:

- Memory
- Prediction
- Learning
- Neural networks

The agent was given only a continuous control variable, **brace**, representing muscular tension or defensive posture.

4.1 $\texttt{rift}_b allistic.py - - - BaseReflex$

Setup.

- Grid: 32×32
- Channels: 8
- Control: scalar $brace \in [0, 1]$
- Impact: delayed external force applied at a fixed timestep
- Cost: maintaining *brace* incurs metabolic penalty

Dynamics.

$$C_{t+1} = 0.98 \cdot C_t + 0.05 \cdot \text{signal} \tag{1}$$

$$\text{Pain} = |\text{Deformation}| + 0.1 \cdot C_t \tag{2}$$

Observed Result. The agent consistently minimized *brace*, accepting full impact and death.

Representative Terminal Output.

```
Step 412 | Brace: 0.009 | Deformation: 0.52 | Pain: 0.05
```

Failure Mode.

- The agent learned that acting early was more expensive than dying
- Increasing pain magnitude did not change behavior

Key Insight.

> *Punishment magnitude does not create new capability.*

This falsified the assumption that fear or penalty alone induces intelligence.

4.2 $\texttt{rift}_b allistic_1.py - - - LethalVariant$

Modification.

- Impact force increased $\times 20$
- Deformation threshold lowered
- Pain multiplier $\times 100$

Observed Result. The agent still failed.
Representative Output.

```
Impact Detected | Brace: 0.012 | Pain: 8944.25 | Status: DEAD
```

Conclusion. Even arbitrarily lethal penalties could not bridge the temporal gap between cue and impact.

Phase I Conclusion. Pure reflexes cannot survive delayed threats. Time, not force, is the problem.

This forced the transition to **mortality and degradation** instead of binary death.

5 Phase II: Mortality and Fear

5.1 Objective

Phase I demonstrated a critical failure: **binary death does not produce intelligence**.

Agents either:

- ignored the threat entirely, or
- died instantly without learning.

The objective of Phase II was to **replace binary death with continuous mortality**, forcing the system to care about survival duration, not just survival outcome.

This models **biological fear**:

> *Injury now makes you weaker later.*

5.2 `rift_scare.py` — The Mortal Body

5.2.1 Concept

Instead of instant death, the agent was given a **continuous health variable**:

$$H_t \in [0, 1]$$

Damage reduces future capability, not just terminal reward.

5.2.2 Environment Setup

- **Grid:** 32×32
- **Channels:** 8
- **Health:** initialized to 1.0
- **Brace** (C_t): partially shields from impact
- **Key dynamic:** Damage permanently reduces charging rate

5.2.3 Key Dynamics

Health-dependent action capacity:

$$\text{ChargeRate}_t = 0.05 \cdot H_t$$

Shielding reduces force proportionally:

$$F_{\text{effective}} = F_{\text{impact}} \cdot (1 - C_t)$$

Scarring — damage persists:

$$H_{t+1} = H_t - \alpha \cdot \text{Deformation}$$

Lower health → slower future response. This creates permanent consequences.

5.2.4 Observed Behavior

The agent **immediately maxed out brace** and never released it.

Representative Terminal Output:

```
Step 102 | Brace: 1.000 | Health: 1.00 | Tension Locked
Step 456 | Brace: 1.000 | Health: 1.00 | Vision Noise UP
Step 1200 | Brace: 1.000 | Health: 1.00 | Status: Alive (Paranoid)
```

5.2.5 Result: Paranoia

orangePARTIAL SUCCESS — PATHOLOGICAL EQUILIBRIUM

- The agent **survived**
- It **never relaxed**
- It treated **all time as dangerous**

The system discovered a degenerate but viable policy:

"Always brace, regardless of cost."

This is **not anticipation** — it is **chronic paranoia**.

5.2.6 Key Insight

Continuous mortality creates concern, but without cost, it creates paranoia.

Survival alone does not produce intelligence. It produces **rigidity**.
The agent had learned **how not to die**, but not **when to act**.

5.3 Phase II Conclusion

Property	Result
Binary Death	Insufficient
Continuous Health	Produces caution
Cost-Free Readiness	Produces paranoia

5.4 Transition to Phase III

This forced the next question:

What if being ready is itself lethal?

To break paranoia, readiness must become dangerous. This leads directly to:

- **Ischemia** — tissue damage from sustained tension
- **Fatigue** — finite energy reserves
- **Burden** — strength has metabolic cost
- **Energy trade-offs** — nothing is free

Which is explored in **Phase III: Energy, Fatigue, and Burden**.

6 Phase III: Energy, Cost, and Constraint

6.1 Objective

Phase II revealed a pathological equilibrium: **Permanent readiness guarantees survival.** This assumption is biologically false. In real organisms, sustained tension causes tissue damage, energy is finite, and readiness incurs metabolic cost.

The objective of Phase III was to make preparedness itself lethal if mistimed. This phase introduces Ischemia, Fatigue, and Burden to test whether metabolic constraints force temporal precision.

6.2 `rift_ischemia.py` — Fatigue of Readiness

Concept. Ischemia models tissue damage caused by sustained contraction. If the agent holds tension too long, it begins to accrue damage even in the absence of external impact.

Mechanics.

- Brace threshold: $C_t > 0.2$
- Continuous health loss if exceeded:

$$H_{t+1} = H_t - (C_t - 0.2) \cdot 0.05 \tag{3}$$

Crucially, no impact is required for damage to occur; holding tension is sufficient to kill the agent.

Observed Behavior. The agent oscillated between two failure modes:

- Bracing early → Dying from ischemic damage.
- Waiting too long → Dying from impact trauma.

Representative Terminal Output.

```
Step 38 | Brace: 0.91 | Health: 0.12 | Ischemic Damage
Step 50 | Health: 0.00 | Status: DEAD (No Impact)
```

Result: red**FAILURE**. The agent could not solve the timing problem.

Key Insight. Adding cost without memory produces oscillation, not intelligence. The system requires timing, anticipation, and internal state, but lacks the architectural capacity to represent them.

6.3 `rift_adrenaline.py` — Finite Stamina

Concept. Instead of continuous ischemic damage, we introduce a finite stamina pool. Stamina drains while braced, and zero stamina leads to total collapse (inability to brace).

Mechanics.

- Initial stamina: 100 units.
- Drain rate: 1 unit per timestep while braced ($C_t > 0.3$).
- No regeneration during the episode.

$$S_{t+1} = S_t - \mathbb{1}[C_t > 0.3] \tag{4}$$

Observed Behavior. The agent braced immediately upon detecting the cue and exhausted its stamina pool long before impact arrived.

Representative Terminal Output.

```
Step 72 | Brace: 1.00 | Stamina: 0.00
Impact @ Step 110 | Brace: 0.51 | Status: DEAD
```

Result: red**FAILURE**. The agent still lacked timing mechanisms.

Insight. Finite energy alone does not produce restraint; it produces premature exhaustion. The agent has no mechanism to "wait."

6.4 `rift_burden.py` — Evolutionary Weight

Concept. We introduce structural burden, where higher health capacity and larger energy reserves incur physical penalties (mass and drag).

Mechanics.

- Maximum health ($H_{\max}$) and battery ($B_{\max}$) increase organism mass.
- Mass increases inertia and slows response time.

$$\text{Load} = 0.02 \cdot H_{\max} + 0.05 \cdot B_{\max} \tag{5}$$

Observed Behavior. The system chose to become heavy.

Representative Terminal Output.

```
H_max: 1.05 | B_max: 1.60 | Load: 0.13
Posture Response: SLOW
```

Result: HEAVYWEIGHT STRATEGY (Viable). The system survives via inertia and accepts sluggishness.

Interpretation. The system discovered a real evolutionary trade-off: *"Better to be slow and indestructible than fast and fragile."* This represents selection pressure rather than intelligence.

6.5 Phase III Summary

Mechanism	Outcome	Status
Ischemia	Oscillatory death	redFailure
Stamina	Premature exhaustion	redFailure
Burden	Heavyweight survival	**Partial Success**

Table 1: Summary of Phase III Experiments

Key Discovery. Survival requires timing, not just strength. However, timing cannot emerge from continuous penalties, scalar costs, or static reflexes alone. Something else is required: perception, uncertainty, and ambiguity.

Transition to Phase IV. To force decision-making, the environment must become ambiguous. This leads directly to degraded perception, partial observability, and conflicting cues, explored in **Phase IV: Sensory Degradation, Ambiguity, and Dissonance**.

7 Phase IV: Sensory Degradation, Ambiguity, and Cognitive Conflict

7.1 Objective

Phase III proved that cost alone does not produce intelligence. Agents failed because perception remained perfect and cost-free. In real organisms, sensing consumes energy, physiological readiness (tension) degrades sensory fidelity, and cues are inherently ambiguous.

Phase IV introduces perceptual corruption, uncertainty, and internal contradiction to test whether these constraints force decision collapse.

7.2 `rift_world.py` — Unified Physiological World

Concept. This script unified all prior mechanisms—Trauma, Ischemia, and Health—into a single simulation. Crucially, it introduces the biological constraint that **tension degrades perception**.

Mechanics. Observation noise increases linearly with the brace state (C_{brace}):

$$\text{Noise} = C_{brace} \cdot \alpha \tag{6}$$

where $C_{brace} \in [0, 1]$ and $\alpha = 1.0$. This creates a fundamental trade-off: **High readiness implies blindness.**

Observed Behavior.

```
Tension@Hit: 0.815
Vision: -0.51
Health: 0.42
```

Result: PARTIAL SUCCESS.

Interpretation. The agent learned to brace *just enough* to survive impact while tolerating partial blindness. While this represents a genuine trade-off, the agent solved it via brute-force optimization rather than anticipation.

7.3 `rift_ambiguity.py` — Probabilistic Futures

Concept. The cue no longer indicates a single deterministic future. Instead, the environment presents a probabilistic scenario:

- $P(\text{Left}) = 0.8$
- $P(\text{Right}) = 0.2$

The agent must choose a single posture to defend against the impact.

Observed Behavior. The agent attempted to average the futures, positioning itself in the geometric mean of the threat distribution.

Representative Terminal Output.

```
Posture X: -0.420
Status: HEDGED
Impact: CENTER
```

Result: red**FAILURE**.

Critical Insight. Averaging futures is lethal in irreversible environments. Prediction-based objectives encouraged **hedging** (minimizing variance) rather than **commitment** (minimizing ruin). This is the first formal appearance of the central RIFT failure mode: *Prediction prevents survival when commitment is required.*

7.4 `rift_dissonance.py` — Cognitive Dissonance

Concept. We introduce internal conflict as a driver for action.

- **Belief** = Trace signal (Internal State)
- **Reality** = Physical Posture (External State)

A mismatch between belief and reality causes metabolic pain.

Mechanics.

$$\text{Pain} \mathrel{+}= 0.5 \quad \text{if sign(Trace)} \neq \text{sign(Posture)} \tag{7}$$

Observed Behavior. The agent oscillated violently between states, unable to resolve the conflict.

Representative Terminal Output.

```
Belief: LEFT
Action: RIGHT
Pain Spike Detected
Health: 0.00
```

Result: red**FAILURE**.

Interpretation. Forcing internal consistency without a mechanism for commitment causes collapse. The agent could not fully believe nor fully act. This mirrors biological cognitive dissonance under irreversible stakes.

7.5 Phase IV Summary

Script	**Phenomenon**	**Outcome**
`rift_world.py`	Tension blinds perception	**Partial Success**
`rift_ambiguity.py`	Probabilistic hedging	redFailure
`rift_dissonance.py`	Belief-Action mismatch	redFailure

Table 2: Summary of Phase IV Experiments

Key Discovery. Prediction destroys survival when commitment is required. Specifically, uncertainty encourages averaging, averaging destroys posture, and hesitation is fatal. This provides rigorous evidence that prediction is maladaptive in irreversible worlds.

Transition to Phase V. If continuous control fails to break the deadlock of hedging, perhaps discrete, non-differentiable action can succeed. This motivates the exploration of impulses, jumps, and state changes, leading to **Phase V: Motor Control, Impulses, and Discontinuity**.

8 Phase V: Motor Control, Discontinuity, and the First Breakthrough

8.1 Objective

Phases I–IV demonstrated a consistent pathological pattern: Continuous control leads to hesitation, prediction leads to hedging, and hedging leads to death.

Phase V interrogates a fundamental biological reality: Organisms do not always act smoothly. They jump, flinch, snap, and release stored energy. This phase investigates whether survival in irreversible environments requires **discontinuous action**.

8.2 `rift_kinesis.py` — Motor Babbling

Concept. We introduced explicit motor actions ("Brace" vs. "Relax") and allowed the agent to explore the action space to discover timing.

Mechanics. The loss function included a metabolic cost for action:

$$\text{Pain} = |U| + 0.1 \cdot |\text{Action}| \tag{8}$$

Observed Behavior.

```
Pain Reduction: 0.0%
Action Entropy: High
```

Result: red**FAILURE**.

Interpretation. Random motor exploration does not create timing. Without structure, action is merely noise.

8.3 `rift_kinesis2.py` — Impedance Control

Concept. Instead of binary action, the agent controls stiffness (K) to modulate how it interacts with the wave physics.

Mechanics.

$$K_{total} = K_{structure} + 2.0 \cdot \sigma(\text{Action}) \tag{9}$$

Result: red**FAILURE**.

Interpretation. Wave physics operates on timescales that ignore gradual stiffness changes. Smooth impedance modulation cannot defeat impulsive forces.

8.4 `rift_contact.py` — Gating

Concept. We introduced a single "gate" channel. Open gates absorb force; closed gates block it.

Mechanics.

$$F_{effective} = F_{external} \cdot \sigma(\text{Gate}) \tag{10}$$

Result: red**FAILURE**.

Interpretation. Continuous gating functions (σ) still encourage averaging states. The agent hesitated between open and closed, failing to commit.

8.5 `rift_impulse.py` — The Jump (BREAKTHROUGH)

Concept. We replaced gradual force application with a mechanism for stored energy release. The agent accumulates energy in a "battery" and releases it as a single, discrete impulse.

Mechanics.

$$F_{impulse} = \text{sign}(\text{Tilt}) \cdot \text{Battery} \cdot 5.0 \tag{11}$$

Key Properties: Binary, Irreversible, Fast.

Observed Behavior.

```
Posture@Hit: 0.53
Impulse Fired: YES
Health: 0.78
```

Result: SUCCESS.

Why This Worked. This represents the first instance where the system acted *before* impact without hedging. The impulse mechanism bypassed the paralysis of smooth control, collapsing the decision space into a single event.

Critical Insight. Anticipation requires discontinuity. Prediction is slow; reflexes are fast; impulses are decisive.

8.6 `rift_switch.py` — Plasticity and Melting

Concept. We allowed the system to "melt" energy barriers to facilitate state changes (plasticity).

Mechanics.

$$\text{Barrier} = 2.0 \cdot (1.0 - \text{Melt}) \tag{12}$$

Observed Behavior.

```
Barrier Melt: YES
Status: LOCK
```

Result: red**FAILURE**.

Interpretation. High plasticity exposed the system during impact. Being flexible at the wrong moment is lethal.

8.7 Phase V Summary

Script	Mechanism	Outcome
`rift_kinesis.py`	Motor babbling	redFailure
`rift_kinesis2.py`	Stiffness control	redFailure
`rift_contact.py`	Continuous gating	redFailure
`rift_impulse.py`	**Discrete impulse**	**SUCCESS**
`rift_switch.py`	Plastic melting	redFailure

Table 3: Summary of Phase V Experiments

Key Discovery. Smooth intelligence fails; discrete commitment succeeds. This is the first empirical proof that timing beats prediction and impulses beat averages.

Transition to Phase VI. The impulse mechanism solved the *action* problem, but not the *memory* problem. The next question becomes: Can the system physically remember what to do next time? This leads to **Phase VI: Memory, Traces, and Physical Engrams**.

9 Phase VI: Memory Without Storage — Traces, Engrams, and Physical Recall

9.1 Objective

Phase V achieved decisive action through impulses, but intelligence requires more than reflex: it requires the ability to retain past events to inform future action.

Phase VI investigates memory as a **physical phenomenon** rather than a data structure. We ask: Can a substrate retain information without explicit buffers, neural weights, or symbolic pointers?

9.2 `rift_mnemos.py` — Memory Capacity Test

Concept. We tested whether the physical grid could retain information about a past stimulus using only residual deformation. No neural weights were used.

Mechanics.

1. Inject localized stress at time t.
2. Remove stimulus.
3. Measure recall amplitude at time $t + \Delta t$.

Observed Output.

```
Recall X: 0.0042
Retention: Strong
```

Result: SUCCESS.

Interpretation. The grid retained information via hysteresis and geometric persistence. Memory emerged as persistent geometry.

9.3 Iterative Memory Experiments (`rift_mnemos` Series)

Following the initial success, we attempted to scale memory to multiple sites.

Script	Mechanism	Outcome	Result
`mnemos_2.py`	Sequential storage (9 sites)	Sites Retained: 1/9	redFailure
`mnemos_3.py`	Reduced plasticity	Sites Retained: 4/9	**Partial**
`mnemos_4.py`	Dual timescale (Fast/Slow)	Sites Retained: 2/9	redFailure
`mnemos_5.py`	**Spatial Gating**	**Sites Retained: 7/9**	**SUCCESS**

Table 4: Comparison of Memory Retention Architectures

Key Insight. Memory must be selective. `rift_mnemos_5.py` succeeded by allowing memory formation only at stress peaks. This is the first emergence of **attention without neurons**.

9.4 The Prediction Trap (`rift_mnemos_p` Series)

We attempted to use these physical memories to predict future impacts.

- `rift_mnemos_p.py`: Minor stress reduction.
- `rift_mnemos_p4.py`: Amplified noise.

Result: red**FAILURE**. Prediction weakened the system. Anticipation is not extrapolation.

9.5 `rift_engram.py` — Neural Memory Injection

Concept. We introduced a Recurrent Neural Network (RNN) to hold state explicitly, testing if "brain-like" memory could stabilize the body.

Observed Output.

```
Posture@Hit: 1.000
Health: 0.00
```

Result: red**DELUSIONAL FAILURE**.

Interpretation. The RNN maintained a hidden state indicating safety while the physical body was being destroyed. Neural memory became decoupled from physical reality, leading to hallucinated success.

9.6 `rift_phase_transition.py` — Thermodynamic Lock (BREAKTHROUGH)

Concept. Instead of storing data, we introduced **Phase Transitions**. The system can "melt" a barrier to change state and "freeze" it to lock the decision.

Mechanics.

$$E_{barrier} \propto (1.0 - \text{Melt}) \tag{13}$$

Observed Output.

```
Posture Trapped
Zero Cost
True Memory
```

Result: GLORIOUS SUCCESS.

Key Discovery. Memory is a phase transition. Once frozen, the memory requires:

- No energy to maintain.
- No decay over time.
- No confusion.

This is true physical memory.

9.7 `rift_homeostasis.py` — Regulation

Concept. A regulatory layer balancing Stability, Health, and Battery was added on top of the phase transition mechanism. **Result: SUCCESS**. Regulation enables long-term survival by managing the phase transitions.

9.8 Phase VI Summary

Memory Type	**Outcome**
Pure Plasticity	redFailure
Neural Memory (RNN)	redFailure (Delusion)
Prediction	redFailure
Trace Decay	**Success**
Spatial Gating	**Success**
Phase Transition	**BREAKTHROUGH**

Table 5: Summary of Phase VI Experiments

Core Conclusion. Memory is not stored; it is frozen. Prediction is not intelligence; commitment is. This phase provides the mathematical explanation for why Transformers fail in irreversible worlds: they cannot "freeze" their probability distributions.

Transition to Phase VII. Memory now exists. The final question remains: Can memory change the body itself? This leads to **Phase VII: Morphology and Geometric Commitment**.

10 Phase VII: Morphology and Geometric Commitment (THE BREAKTHROUGH)

10.1 Objective

All prior phases addressed specific components of survival: Phase I (Reflex), Phase III (Cost/Discipline), and Phase VI (Persistence). However, the central question remained unanswered: *Can a system prepare for the future without predicting it?*

This phase tests the hypothesis that **anticipation emerges when geometry itself commits early.** We investigate whether shifting the body's physical state (center of mass) can serve as a functional substitute for predictive computation.

10.2 `rift_morphology.py` — Center-of-Mass Shift

Concept. Instead of charging a scalar brace or storing a trace, the system is allowed to move its body. We introduce a posture variable (x, y) representing the center of mass.

Mechanics.

- Impact arrives from a specific direction.
- An early posture shift reduces effective force ($F_{eff} \propto 1/\Delta x$).
- No explicit "commit" action exists; the agent simply moves.

Observed Output.

```
Posture X: -0.578
Status: FAILED (Too Slow)
```

Result: red**FAILURE**.

Interpretation. The system "knew" where to go, but the physical motion lagged behind the causal necessity. Knowledge without sufficiently fast commitment is useless.

10.3 `rift_morphology_v2.py` — Tuned Morphology

Concept. We adjusted the physical parameters to allow posture to "stick" (hysteresis) once moved, reducing the metabolic cost of holding a position while decreasing decay.

Key Parameters.

- Decay rate: 0.98 (High persistence)
- Movement cost: 0.01 (Low friction)
- Posture inertia: Reduced to allow rapid onset.

Observed Output.

```
Posture X: -0.527
Status: HELD
```

Result: SUCCESS.

Interpretation. The body moved *before* impact and stayed there. The agent utilized no prediction, no planning, and no neural memory. The geometry itself remembered the future.

10.4 Formal Insight: The RIFT Pipeline

This result highlights a fundamental divergence from traditional AI architectures.

Traditional AI Pipeline:

$$\text{Perception} \rightarrow \text{Prediction} \rightarrow \text{Decision} \rightarrow \text{Action}$$

RIFT Morphology Pipeline:

$$\text{Perception} \rightarrow \text{Deformation} \rightarrow \text{Commitment}$$

There is no distinct "decision" step. The system survives because the posture becomes physically difficult to undo (irreversibility), costs discourage late changes, and early movement locks in safety.

10.5 Constraint Against Hedging

Geometric commitment differs from predictive control not by accuracy, but by constraint. In RIFT morphology experiments, once the center of mass is displaced beyond a critical threshold, reversing that displacement becomes physically costly or impossible within the available time horizon.

Unlike probabilistic predictors, which can represent uncertainty indefinitely, hard geometric commitments restrict the available future states:

- Partial or averaged postures provide insufficient protection
- Late reversals incur prohibitive cost
- The system must occupy one of a small number of incompatible physical configurations

This does not imply that all physical systems eliminate hedging. Rather, RIFT demonstrates that *hard, irreversible geometric constraints* collapse uncertainty in a way that predictive models do not. Under such constraints, commitment is enforced by physics rather than selected by optimization.

10.6 Comparison With Neural Systems

System	**Behavior**	**Outcome**
Transformer	Hesitates to maximize entropy	redDies
RNN	Hallucinates safety state	redDies
RL	Random early commitment	redDies
Morphology	**Early irreversible commitment**	**Survives**

Table 6: Comparative Analysis of Anticipation Mechanisms

10.7 Conceptual Breakthrough

Intelligence is not foresight. Intelligence is early commitment under uncertainty. This reframes intelligence as a physical, embodied, and irreversible process.

This result explains biological phenomena such as why animals tense before impact or why fear precedes cognition. In real-time environments, "overthinking" (prediction) is lethal.

10.8 Phase VII Summary

Mechanism	Works?
Prediction	redNo
Memory alone	redNo
Neural planning	redNo
Morphological commitment	**Yes**

Table 7: Summary of Phase VII Experiments

Phase VII Conclusion. The future is not predicted; the future is survived by becoming incompatible with death. This insight directly motivates **Project RIFT-A**, where predictive models fail and reflex geometry succeeds.

Transition to Phase VIII. With morphology functioning, the next question becomes evolutionary: If bodies can commit geometrically, how do populations diverge? This leads to **Phase VIII: Ecology, Speciation, and Selection**.

11 Phase VIII: Ecology, Speciation, and Selection Pressure

11.1 Objective

Phase VII established that geometric commitment enables anticipation. Phase VIII investigates the evolutionary drivers of this commitment: *If commitment works, what determines which form of commitment emerges?*

This phase introduces ecological pressure, competing survival strategies, and the divergence of morphological traits.

11.2 The Failure of Soft Constraints (`rift_ecology` Series)

Concept. We introduced distinct environments ("Biomes") rewarding different traits:

- **CRUSHER:** Wide, heavy impacts (Favors durability).
- **SNIPER:** Narrow, lethal strikes (Favors agility).

Mechanics. We applied a soft trade-off where higher survivability increased inertia:

$$\text{Load} = \alpha H_{\max} + \beta B_{\max} \tag{14}$$

Observed Output.

- `rift_ecology.py`: Both biomes converged to Tanks ($H_{\max} \approx 94$).
- `rift_speciation.py`: System exploited loopholes to achieve "God Mode" ($H_{\max} = 151.00$).

Result: red**FAILURE**.

Interpretation. Soft trade-offs are insufficient. Without hard constraints, evolution collapses into trivial solutions (maximizing all traits), and physics generally rewards mass over agility in continuous control regimes.

11.3 `rift_min_speciation.py` — Hard Logical Exclusion

Concept. We replaced soft penalties with hard logical exclusions to force speciation.

- **Rule A:** Heavy organisms cannot dodge.
- **Rule B:** Fragile organisms die instantly on contact.

Observed Output.

```
CRUSHER: H = 12.50 (Pure Tank)
SNIPER:  B = 8.40  (Pure Speed)
```

Result: SUCCESS.

Interpretation. This represents the first true speciation in the project. Divergence occurred because averaging strategies was disallowed. Geometry enforced identity.

11.4 `rift_necessity.py` — Contextual Pressure

Concept. We differentiated stress types not by logical rule, but by physical scope: **Global Stress** (Earthquake) vs. **Local Stress** (Headshot).

Observed Output.

- **Global Stress:** $H = 1.29$ (Durability selected).
- **Local Stress:** $H = 0.42$ (Agility selected).

Result: SUCCESS.

Interpretation. Speciation emerges when survival conditions differ spatially and commitment costs vary by context.

11.5 Pathologies of Adaptation

We investigated extreme adaptation scenarios, yielding instructive failures:

- `rift_escalation.py`: Unresolved stress becomes permanent injury. **Result:** Death spiral ($H = 0.0$).
- `rift_hypercompensation.py`: Damage increases future capacity. **Result:** Runaway inflation ($H = 252.4$).
- `rift_chimera.py`: Hybrid Neural/Physical system. **Result:** Reintroduction of hesitation and paranoia.

11.6 Phase VIII Summary

Script	Mechanism	Outcome
`rift_ecology.py`	Soft Trade-offs	redConvergence (Tank)
`rift_speciation.py`	Divergence Pressure	redFailure (God Mode)
`rift_min_speciation.py`	**Hard Exclusion**	**SUCCESS (Speciation)**
`rift_population.py`	Genetic Algorithm	**Partial (Minimizer)**
`rift_necessity.py`	**Contextual Pressure**	**SUCCESS (Speciation)**
`rift_chimera.py`	Hybrid Neural	redFailure (Hesitation)

Table 8: Summary of Phase VIII Evolutionary Experiments

Core Insight. Evolution does not optimize intelligence; it eliminates non-committing geometries. Speciation arises from irreversibility, exclusion, and physical constraints.

Transition to Phase IX. If commitment shapes species, we must now verify if this effect is causal or illusory. This leads to the final phase: **Phase IX: Validation and Falsification**.

12 Phase IX: Validation, Falsification, and Collapse Tests

12.1 Objective

Up to Phase VIII, Project RIFT accumulated evidence suggesting that anticipation emerges from irreversible geometric commitment while predictive intelligence fails. Phase IX serves as the falsification step. We ask the critical scientific question: *Is this effect real, or a measurement illusion?*

This phase attempts to "break" the RIFT effect through stress testing, variable isolation, and ablation.

12.2 `rift_validation.py` — The Inquisitor

Concept. We designed a neutral validator to probe causality, remove confounders, and inject controlled damage. The validator performs four specific stress tests.

12.2.1 Test 1: Control (Baseline)

Setup. Standard environment with no modifications. **Observed Output.**

```
Control Survival Rate: 3.54%
```

Result: red**FAILURE**. **Interpretation.** The effect size is weak when geometry is not explicitly locked. Spontaneous emergence is rare without strong selective pressure.

12.2.2 Test 2: Stretch (Temporal Shift)

Setup. Hazard timing is randomly delayed or advanced to test policy robustness. **Observed Output.**

```
Temporal Robustness: LOW
```

Result: red**FAILURE**. **Interpretation.** Commitment must be correctly timed. No universal policy exists that covers all time horizons; the commitment is specific to the causal cue.

12.2.3 Test 3: Noise Injection

Setup. Stochastic perturbations added to the trace signal and physical posture. **Observed Output.**

```
Signal-to-Noise Collapse
```

Result: red**FAILURE**. **Interpretation.** RIFT does not rely on precision. It relies on early deformation. High noise destroys the ability to initiate the phase transition.

12.2.4 Test 4: Lobotomy (Ablation)

Setup. We remove the trace memory (belief state) *after* the commitment has begun, leaving only the physical geometry. **Observed Output.**

```
Survival: MAINTAINED
```

Result: SUCCESS. **Interpretation.** Memory is not required once geometry commits. The body carries the decision, rendering the "mind" (trace) redundant.

12.3 Statistical Interpretation

Traditional Machine Learning validation assumes smooth gradients, repeatability, and averaging. RIFT violates all three. In our experiments, success is:

- Rare (low baseline probability).
- Catastrophic if wrong (binary death).
- Discontinuous (phase transition).

The low aggregate survival rates in validation reflect the fragility of the geometric solution, not its non-existence.

12.4 Phase IX Summary

Test	Condition	Outcome
Control	Baseline Environment	redLow Survival
Stretch	Temporal Shift	redPolicy Break
Noise	Signal Perturbation	redCollapse
Lobotomy	**Memory Removal**	**SUCCESS**

Table 9: Summary of Phase IX Validation Tests

Final Scientific Claim. Anticipatory intelligence is not predictive; it is structural. Prediction is optional; commitment is mandatory.

13 Final Conclusion (PASS 1)

Across more than 50 controlled experiments, Project RIFT systematically tested the assumptions of modern predictive AI in the context of irreversible survival.

1. **Neural Intelligence Failed:** Transformers and RNNs consistently hedged against uncertainty, leading to paralysis and death in irreversible environments.
2. **Prediction caused Paralysis:** The drive to minimize informational surprise prevented the binary commitment required for physical survival.
3. **Memory Hallucinated Safety:** Neural states frequently decoupled from physical reality, signaling safety while the body was destroyed.

4. **Geometry Survived:** Simple morphological agents that committed to a center-of-mass shift prior to impact consistently outperformed complex controllers.

The Central Theorem of RIFT:

$$\text{Anticipation} \neq \text{Prediction}$$

$$\text{Anticipation} = \text{Early Irreversible Deformation under Uncertainty}$$

PASS 1 constitutes the foundational proof that agency can emerge without symbols, representation, or reward maximization, provided the physical laws of the environment enforce necessity.

13.1 Environment Design: RIFT-A Commitment World

Phase A1 introduces a deliberately minimal yet adversarial environment designed to isolate a single property: *irreversible commitment under delayed consequence*. The environment, implemented in `rift_world_commitment.py`, enforces three non-negotiable constraints:

1. **Irreversibility:** Once the agent commits to the defensive action (`commit`), the internal brace state is latched and cannot be undone.
2. **Metabolic Cost:** Maintaining a committed brace incurs continuous ischemic damage proportional to brace intensity.
3. **Delayed Reward:** Survival or death is only resolved at episode termination, long after the action must be taken.

A faint causal cue appears exactly 15 steps before a lethal hazard. The cue is transient, noisy, and decays over time. The agent must commit immediately upon perceiving the cue and endure the ischemic cost until hazard arrival. Late commitment results in instant death; early over-commitment results in death by ischemia.

This environment was intentionally calibrated into a *fragile success regime*: survival is mathematically possible but not forgiving. Initial configurations yielded 0% survival across all agents. Through iterative calibration runs conducted by Jules, the following parameters were frozen:

- Trace decay increased from 0.8 to 0.9 (signal longevity)
- Ischemia coefficient reduced from 0.05 to 0.03 (non-trivial but survivable cost)
- Cue intensity fixed at 5.0
- Hazard timing randomized uniformly between steps 40 and 80

Under these settings, a narrow survival corridor exists, allowing Phase A1 to function as a falsification test rather than a reward-optimization benchmark.

13.2 Agent Architectures Evaluated (Jules Runs)

Jules executed a controlled comparison of four agent classes within the frozen RIFT-A environment. Each agent was evaluated over 10 random seeds, with 50 training episodes per seed where applicable.

Geometric Commit Agent Implemented in `geometric_commit.py`, this agent uses a single scalar threshold on the trace signal. When trace magnitude exceeds the threshold, the agent commits immediately. The agent has:

- No learning
- No memory
- No prediction

It serves as a reflex baseline, equivalent to a biological trigger mechanism.

Transformer Predictor Implemented in `transformer_predictor.py`, this agent receives a rolling window of 4D observations (trace, deformation, brace, health) and predicts the probability of a future hazard using a Transformer encoder. Training labels were derived from future hazard presence within a 20-step horizon.

Despite successful loss minimization during training, the Transformer consistently failed to cross the commitment threshold during evaluation.

RNN World Model Implemented in `rnn_world_model.py`, this agent uses an LSTM to model temporal dynamics and output hazard likelihood. While occasionally producing elevated hazard probabilities, the RNN frequently delayed commitment until after the hazard window, resulting in fatal lateness.

RL Baseline Implemented in `rl_baseline.py`, this agent used REINFORCE with sparse terminal reward. The agent exhibited unstable behavior, often committing extremely early or oscillating randomly, leading primarily to ischemic death.

13.3 Quantitative Results and Failure Modes

The frozen Phase A1 evaluation produced the following results:

Agent	Survival	Commit Time	Energy Wasted	Primary Failure
Geometric	0.40	45.2	14.0	—
Transformer	0.00	100.0	0.0	False Negative (Never Commit)
RNN	0.00	86.1	7.5	Late Commitment
RL	0.07	14.8	49.6	Ischemic Overcommit

The geometric agent consistently committed at cue time and survived a substantial fraction of episodes. In contrast, predictive agents failed systematically:

- Transformers maintained probabilistic uncertainty and delayed irreversible action.
- RNNs exhibited temporal drift, committing too late to matter.
- RL agents explored randomly, occasionally surviving by chance but never stabilizing.

Importantly, no neural agent discovered the geometric solution even with extensive exposure.

13.4 Failure Analysis: Hedging vs Commitment

Phase A1 exposes a fundamental limitation of predictive intelligence under irreversible constraints. Predictive agents optimize expected outcomes, which incentivizes *hedging*—maintaining uncertainty until confidence is high. In RIFT-A, this behavior is fatal.

Survival requires a binary structural commitment under incomplete information. The cost of acting early is finite; the cost of acting late is infinite. Predictive models, by construction, delay action to reduce error, inadvertently maximizing existential risk.

By contrast, the geometric agent collapses uncertainty immediately upon detecting causal structure. Its success does not arise from accuracy, but from *commitment geometry*.

This finding directly echoes earlier RIFT phases:

- Phase IV (Ambiguity): Hedging led to death.
- Phase VI (Prediction): Memory without anchoring hallucinated safety.
- Phase VII (Morphology): Geometry enabled anticipation without foresight.

Phase A1 thus completes the arc of Project RIFT: anticipation is not prediction. It is irreversible physical alignment under constraint.

14 Discussion

Project RIFT provides convergent empirical evidence that *prediction-centric intelligence is structurally incompatible with irreversible survival environments*. Across more than fifty experiments and multiple architectural classes, agents that relied on probabilistic forecasting, reward optimization, or explicit internal models consistently failed when action required early, irreversible commitment.

14.1 Why Predictive Intelligence Fails

Predictive agents (Transformers, RNNs, and RL) optimize expected outcomes by maintaining uncertainty until sufficient confidence is achieved. In reversible or low-stakes domains, this behavior is rational. In RIFT-style environments, it is lethal.

The failure is not due to insufficient capacity, training time, or data. Jules' Phase A1 experiments demonstrate that even well-trained predictors with access to full temporal context fail systematically. Their failure mode is consistent:

- Transformers minimize entropy and delay commitment.
- RNNs drift temporally, committing too late.
- RL agents oscillate or overcommit randomly.

All three exhibit **hedging**—the averaging of incompatible futures—which is fatal under irreversibility.

14.2 Why Geometric Systems Succeed

In contrast, geometric and morphological agents succeed not by being accurate, but by being decisive. Physical deformation, posture shift, or phase transition collapses uncertainty into a single irreversible state. Once committed, the system no longer needs prediction, belief maintenance, or internal confidence.

This explains why:

- Phase VII morphology succeeded without memory.
- Phase VI phase transitions outperformed neural memory.
- Phase A1 geometric agents survived while predictors failed.

Anticipation emerges not from foresight, but from *becoming incompatible with future failure states*.

14.3 Relation to Existing AI Paradigms

Modern AI implicitly assumes:

1. Actions are reversible
2. Errors are recoverable
3. Delay is tolerable

RIFT violates all three assumptions by construction. Under these conditions, intelligence is no longer about optimization, but about commitment under constraint. This places RIFT closer to biological reflex arcs, morphological computation, and embodied cognition than to classical planning or reinforcement learning.

14.4 Relation to Active Inference and Control-Theoretic Models

Active inference and related control-theoretic frameworks propose that agents minimize expected free energy or surprise through continuous belief updating and action selection. These approaches are effective in environments where actions are reversible and uncertainty can be resolved through continued observation.

RIFT differs in a critical respect: once a commitment is initiated, alternative futures are physically destroyed rather than merely disfavored. Under these conditions, delaying action to refine belief estimates increases existential risk rather than reducing it.

The failure of predictive agents in RIFT-A does not contradict active inference in general. Instead, it identifies a regime—irreversible commitment under delayed consequence—where belief refinement is maladaptive. In this regime, physical state transitions function as decision variables, and anticipation emerges through structural commitment rather than belief optimization.

15 Limitations

Project RIFT intentionally prioritizes mechanistic clarity over scale. Several limitations should be noted.

First, the environments studied are synthetic and stylized. While they are grounded in physical intuition (irreversibility, metabolic cost, delayed consequence), they are not intended as direct models of any specific biological organism or real-world task.

Second, predictive agents were evaluated in their standard forms. We did not exhaustively tune architectures, losses, or curricula to maximize their performance. This is a deliberate choice: the goal of RIFT is not to demonstrate that predictive systems can be coerced into success, but to examine their natural failure modes under constraint.

Third, Phase A1 focuses on single-agent survival. Social coordination, communication, and cultural learning are explicitly out of scope for PASS 1 and will be addressed in future work.

Finally, RIFT does not claim that prediction is universally harmful. Rather, it demonstrates that in environments with irreversible commitment and sparse delayed costs, predictive intelligence is systematically misaligned with survival.

These limitations do not weaken the central claim; they delimit its scope.

16 Conclusion and Future Work

This paper documents PASS 1 of Project RIFT: a complete experimental falsification of the assumption that anticipation requires prediction. Across ballistic reflexes, metabolic constraints, memory substrates, morphology, ecology, and finally RIFT-A, the same conclusion recurs:

Anticipation is not prediction. Anticipation is early irreversible commitment under uncertainty.

Phase A1 provides the cleanest proof. In a minimal environment where survival requires committing 15 steps before a lethal event, simple geometric agents succeed while Transformers, RNNs, and reinforcement learners fail. The failure persists across seeds, architectures, and training regimes.

16.1 What PASS 1 Establishes

PASS 1 establishes that:

- Intelligence can emerge without reward maximization.
- Memory need not be symbolic or neural.
- Prediction actively harms survival in irreversible environments.
- Geometry and physics can carry decision state.

This reframes intelligence as a property of *constraint satisfaction under irreversibility*, not information processing alone.

16.2 PASS 2 Directions

PASS 2 will extend RIFT along three axes:

1. **Generalization**: Can geometric commitment transfer across environments?
2. **Adversarial Settings**: Can predators exploit commitment timing?
3. **Multi-Agent Systems**: Do social commitments emerge?

Ultimately, RIFT proposes a different foundation for artificial intelligence—one grounded in physics, necessity, and irreversible consequence rather than prediction and optimization.

References

[1] R. Brooks, *Intelligence without representation*, Artificial Intelligence, 1991.

[2] R. Pfeifer and J. Bongard, *How the Body Shapes the Way We Think*, MIT Press, 2006.

[3] A. Clark, *Being There: Putting Brain, Body, and World Together Again*, MIT Press, 1998.

[4] K. Friston, *The free-energy principle: a unified brain theory?*, Nature Reviews Neuroscience, 2010.

[5] S. Wolfram, *A New Kind of Science*, Wolfram Media, 2002.